

ECE496 Final Report

DISTRIBUTED MINIMAX

Project No. 307

Authors:

Steven BLENKINSOP (123456789)

Jonathan GRAY (998198588)

Supervisor:

Prof. Jorg LIEBEHERR

Administrator:

John TAGLIONE

Submitted March 19, 2015

1 Executive Summary

2 Group Highlights and Individual Contributions

3 Table of Contents

1	Executive Summary	1
2	Group Highlights and Individual Contributions	2
3	Table of Contents	3
4	Background and Motivation	5
5	Project Goal	5
6	Project Requirements	6
6.1	Functional Requirements	6
6.2	Constraints	6
6.3	Objectives	6
7	Final Design	6
7.1	System-Level Overview	7
7.2	Module-Level Descriptions	7
7.2.1	User-Master RPC	7
7.2.2	Job Pool Creation	8
7.2.3	Job Allocation	9
7.2.4	Master-Slave RPC	10
7.2.5	Slave Work	10
7.2.6	Result Consolidation	11
8	Testing and Verification	12
8.1	Requirement: <i>shall correctly implement the minimax algorithm</i>	12
8.2	Requirement: <i>allow for distribution over a number of nodes between 2 and 10</i>	12
8.2.1	Test: Two-Node Cluster	12
8.2.2	Test: Ten-Node Cluster	12
8.3	Requirement: <i>allow the user to set a time limit on the computation . .</i>	14
8.4	Constraint: <i>must be executable on Mac OS 10.9, Windows 8.1, and Ubuntu 14.04</i>	14

8.5	Objective: <i>evaluate more game tree nodes than the non-distributed implementation</i>	14
8.6	Objective: <i>should maximize throughput on different combinations of fast and slow processors</i>	14
9	Summary and Conclusions	14
10	References	15
11	Appendix A: Gantt Chart History	16
12	Appendix B: Original Validation and Acceptance Tests	17
13	Appendix C: Code Samples	18
13.1	Protobuf Specifications	18

4 Background and Motivation

Some of the earliest software written for programmable computers sought to enable them to play abstract strategy games such as chess and checkers [1]. In the late 1940s, artificial intelligence pioneers like Claude Shannon, John von Neumann, and Alan Turing independently developed game-playing algorithms that relied on the concept of minimaxing, an optimal algorithm for game-playing which searches for the most advantageous move assuming the opponent plays perfectly [2].

Parallelization of general-purpose computation on a graphics processing unit (GPU) has been popular for the last 15 years, but GPU frameworks and hardware remain complicated and expensive compared to central processing unit (CPU) cluster technology [3]. Recently, researchers have developed algorithms to parallelize the minimax algorithm on GPUs and on CPU clusters [4][5]. Given the availability and low cost of commodity server hardware and cloud computing platforms, and the relative ease with which commodity-grade processors can be replaced or added to a cluster, CPU clusters are often preferable to GPUs in achieving scalable parallelization of general-purpose computing [6]. For these reasons, we have focused on CPU clusters for this project.

This project aimed to produce a software framework for distributing the execution of the minimax algorithm over an arbitrarily-sized cluster of computers. The framework allows an end-user to describe the game tree and provide heuristics for evaluating the state of a game that has not yet terminated, while abstracting away the mechanics of the distributed nature of the computation. It also provides a simple interface to add, remove, and replace computers participating in the minimax search.

The primary motivation for this project was the group members interest in game artificial intelligence and distributed systems.

5 Project Goal

The goal of this project was to create a software framework with which a user can distribute the evaluation of a user-defined perfect-information zero-sum sequential game using the minimax algorithm over an arbitrary number of computers.

6 Project Requirements

6.1 Functional Requirements

The design shall:

1. correctly implement the minimax algorithm for user-defined perfect-information zero-sum sequential games
2. allow for distributing evaluation over a number of nodes between 2 and 10^1 ; and
3. allow the user to set a time limit (at least 1 second) on the computation

6.2 Constraints

The resulting software must be executable on all of:

1. Mac OS 10.9 or above
2. Windows 8.1 or above
3. Ubuntu 14.04 or above

6.3 Objectives

The resulting software should:

1. evaluate more game states per unit of time compared to an equivalent implementation of the minimax algorithm running on single computer
2. efficiently utilize processors of different speeds

7 Final Design

Section 7.1 provides an overview of the final design and the interactions between the various modules.

Following that, Section 7.2 provides a more detailed description about the major design components and processes described in the system overview.

¹TODO

7.1 System-Level Overview

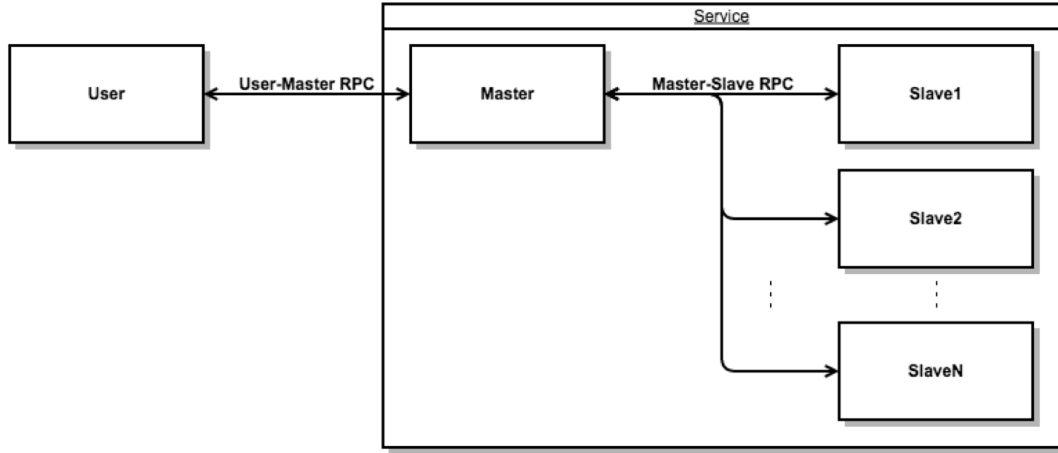


Figure 1: Component diagram for the distributed minimax service showing the separation of *user*, *master*, and *slave* processes

A single *master* process and a number of identical *slave* processes communicate with each other to provide a network *service* to a *user* process. As illustrated in Figure 1, the *user* and the *slaves* communicate only with the *master* (over the *User-Master RPC* and *Master-Slave RPC*, respectively) and not with each other.

On receiving a request from the *user* over the *User-Master RPC*, the *master* creates a *job pool* by breaking down the Minimax algorithm into discrete jobs that can be performed concurrently, and distributes these jobs to the slaves over the *Master-Slave RPC*. The *slaves* execute their jobs and return the results of their analyses to the *master*. The *master* consolidates these results into a single *best move* suggestion and returns this to the *user*.

7.2 Module-Level Descriptions

This section gives a formal description of the major modules referenced in the previous section. Each module's inputs and outputs are listed along with a brief description of the module's purpose. The final design of each module is also described, and areas for potential improvement are noted in some cases.

7.2.1 User-Master RPC

Module Inputs:

- Serialized game state

- Time limit

Module Outputs:

- Suggestion for best move

Module Description:

This module represents network interaction between the *user* and the *master*.

The user makes a request to the master with a serialized game state and a maximum time allowed. The Master node is expected to confer with the slaves and respond within the time limit with its suggestion for a move to be played.

Final Design:

The open-source *Protobuf*² library was used to describe the message formats and generate RPC code. For an exact specification of the RPC, see Section 13.1 in Appendix C.

7.2.2 Job Pool Creation

Module Inputs:

- Game state to be analyzed
- A set of slaves that are available to do work

Module Outputs:

- A set of jobs that can be worked independently and concurrently

Module Description:

This module is a component of the *master*.

Given a game state, this module generates a set of jobs that can be worked independently. The number of jobs created will depend on the number of slaves available to do work. In particular, it should produce at least as many jobs as there are slaves, so no slave is left with no work to do.

²TODO

Final Design:

A *job* in the final design is a game state that is a descendent³ of the game state to be analyzed. This module will examine the direct descendents of the given game state, then each of their direct descendents, etc., until the number of descendents is greater than the number of slaves. When that condition is met, these game states are returned as jobs.

7.2.3 Job Allocation**Module Inputs:**

- A set of jobs that can be worked independently and concurrently
- A set of *slaves* that are available to do work

Module Outputs:

- Exactly one *slave* assigned to each job

Module Description:

This module is a component of the *master*.

Given a set of jobs and a set of *slaves*, this module decides which *slaves* will work on which jobs.

Final Design:

The final design assigns jobs to slaves in a round-robin fashion. For example, if there are 10 jobs and 7 slaves, then the first three slaves will be assigned two jobs each, and the other four slaves will be assigned only one job. This method ensures that the *number* of jobs assigned to each slave is fairly even, but is agnostic about *which* jobs are assigned to which slave.

This is a major area for improvement. Given more time, this module should attempt to measure the relative importance of jobs and the effectiveness of slaves at performing work, and assign more important jobs to more effective slaves.

³The "children" of a game state are the game states that can be reached by making a single move. For example, an empty tic-tac-toe board has nine children, representing a piece played in each of the nine empty spots. The "descendents" of a game state are its children, its children's children, etc.

7.2.4 Master-Slave RPC

Module Inputs:

- A set of jobs
- Time Limit

Module Outputs:

- Result of analysis for each job

Module Description:

This module represents network interaction between the *master* and the *slaves*.

The master delivers to each slave a set of jobs to be performed and a maximum time allowed. The slaves return their results to the Master within the allowed time.

This module deals primarily with serialization and communication with the master node, while the *Slave Work* module below deals primarily with performing analysis.

Final Design:

The open-source *Protobuf*⁴ library was used to describe the message formats and generate RPC code. For an exact specification of the RPC, see Section 13.1 in Appendix C.

7.2.5 Slave Work

Module Inputs:

- A set of jobs
- Time Limit

Module Outputs:

- Result of analysis for each job

⁴TODO

Module Description:

This module is a component of the *slaves*.

This module performs executes the jobs that were allocated by the master. This involves running the Minimax algorithm on subsets of the game tree. The results of the jobs are returned to the *Master-Slave* RPC module to be returned to the master.

Final Design:

This module receives a set of game states to analyze and a time limit, and produces a heuristic evaluation of each game state. For each game state, concurrently, it will run an iterative-deepening Minimax algorithm⁵, analyzing to increasing depths until time runs out. When the time limit has elapsed, the results of the last complete iteration for each game state are returned.

7.2.6 Result Consolidation**Module Inputs:**

- Results of analysis for each job

Module Outputs:

- Suggestion for best move

Module Description:

This module is a component of the *master*.

This module combines the results of the analyses that the slaves performed to produce a single best move suggestion to return to the user.

Final Design:

Given a set of game states that are descendents of the "parent" game state for which the user requested analysis, this module runs the Minimax algorithm rooted at the parent, using the slaves' heuristic evaluations. The output of this Minimax evaluation is returned to the user as the suggested best move.

⁵The Minimax algorithm is run first with a cut-off depth of 1, then 2, then 3, etc. until time runs out.

8 Testing and Verification

This section introduces a test for each requirement, constraint, and objective introduced in Section 6. The results of each test are included along with a short discussion concluding whether or not each requirement was met.

8.1 Requirement: *shall correctly implement the minimax algorithm*

TODO

8.2 Requirement: *allow for distribution over a number of nodes between 2 and 10*

8.2.1 Test: Two-Node Cluster

Test description: Analyze a Connect-4 game state using a cluster of 2 slaves

Pass Criteria: The analysis completes without error, and the result of the analysis is an optimal move.

Test Result: PASS

Included here is the raw log output from the user process (Figure 2) and the master process (Figure 3) when this test was run.

We can verify that this is a 2-slave cluster in line 8 of Figure 3. We can further verify that the analysis completed without error and that the result of the analysis (at the bottom of Figure 2) was an optimal move⁶. Since the analysis completed successfully and without error for a cluster of 2 slaves, this test is considered to have passed.

8.2.2 Test: Ten-Node Cluster

Test description: Analyze a Connect-4 game state using a cluster of 10 slaves

⁶The rules of Connect-4 are not included in this document, but are widely available. In this case, the recommended move is optimal because any other move would lead to an immediate win by X playing in the same spot.

```

1 2015/03/15 23:25:48 asking about state:
2
3 | | | | | | | |
4 | | | | | | | |
5 | | | | | | | |
6 | | | | |0| | |
7 | | | |0|X| | |
8 | | |0|X|X|X| |
9 2015/03/15 23:25:48 making request, time limit 3000 millis
10 2015/03/15 23:25:52 rpc finished in 3.180523374s
11 2015/03/15 23:25:52 recommended move:
12
13 | | | | | | | |
14 | | | | | | | |
15 | | | | | | | |
16 | | | | |0| | |
17 | | | |0|X| | |
18 | | |0|X|X|X|0|

```

Figure 2: The raw output of the user process with a 2-node slave cluster

```

1 2015/03/16 03:25:35 Slave service listening on port 14782
2 2015/03/16 03:25:35 User service listening on port 14783
3 2015/03/16 03:25:39 slave connected with 0 past results
4 2015/03/16 03:25:41 slave connected with 0 past results
5 2015/03/16 03:25:49 DoWork called
6 2015/03/16 03:25:49 expanding rootState for 2 slaves
7 2015/03/16 03:25:49 created 7 jobs from root state
8 2015/03/16 03:25:49 distributing 7 jobs to 2 slaves
9 2015/03/16 03:25:52 slave connected with 3 past results
10 2015/03/16 03:25:52 AddResult {X:1928440315904 Y:208339468288 IsXMove:true}

```

Figure 3: The first 10 lines of raw output of the master process with a 2-slave cluster. The complete log is too long to list here.

Pass Criteria: The analysis completes without error, and the result of the analysis is an optimal move.

Test Result: PASS

A cluster of 10 slaves was used to analyze the same Connect-4 state, and the same (optimal) result was returned. Relevant lines of the master process's log output are shown in Figure 4.

It can be seen in this log output that 10 slaves are used. Since the analysis completed successfully and without error using a cluster of 10 slaves, this test is considered to have passed.

```
1 2015/03/16 02:55:22 DoWork called
2 2015/03/16 02:55:22 expanding rootState for 10 slaves
3 2015/03/16 02:55:22 created 13 jobs from root state
4 2015/03/16 02:55:22 distributing 13 jobs to 10 slaves
```

Figure 4: An excerpt of the raw output of the master process with a 10-slave cluster. The complete log is too long to list here.

8.3 Requirement: *allow the user to set a time limit on the computation*

8.4 Constraint: *must be executable on Mac OS 10.9, Windows 8.1, and Ubuntu 14.04*

8.5 Objective: *evaluate more game tree nodes than the non-distributed implementation*

8.6 Objective: *should maximize throughput on different combinations of fast and slow processors*

9 Summary and Conclusions

10 References

<https://chessprogramming.wikispaces.com/Iterative+Deepening>

11 Appendix A: Gantt Chart History

12 Appendix B: Original Validation and Acceptance Tests

13 Appendix C: Code Samples

This appendix includes raw samples of code from the project to serve as further information for interested developers, and as proof of work for ECE496 administrators.

13.1 Protobuf Specifications

User-Master RPC

```
1 message DoWorkRequest {
2     required bytes state = 1;
3     required uint64 timeLimitMillis = 2;
4 }
5
6 message DoWorkResponse {
7     required bytes move = 1;  // suggestion for next state
8 }
9
10 service UserService {
11     rpc DoWork(DoWorkRequest) returns (DoWorkResponse);
12 }
```

Master-Slave RPC

```
1 message GetWorkRequest {
2     message Result {
3         required bytes state = 1;
4         required int32 value = 2;  // should match Value type in game/game.go
5         required int64 numStatesAnalyzed = 3;
6     }
7     // optionally return the results of a previous workload
8     repeated Result result = 1;
9 }
10
11 message GetWorkResponse {
12     repeated bytes state = 1;
13     required uint64 timeLimitMillis = 2;
14 }
15
16 service SlaveService {
17     rpc GetWork(GetWorkRequest) returns (GetWorkResponse);
18 }
```